

More SQL: NULL, Outer Joins

NULL

(null > 0) = null
 (null + 1) = null
 (null = 0) = null
 (null AND true) = null
 null is null = true

Some truth tables

AND	T	F	NULL
T	T	F	NULL
F	F	F	F
NULL	NULL	F	NULL

OR	T	F	NULL
T	T	T	T
F	T	F	NULL
NULL	T	NULL	NULL

NULL comparisons: unknown

Null is "unknown" or "maybe"

null > 16? Unknown!

left AND right: True if BOTH left and right are true;

NULL AND true? Could be true if NULL was true: = NULL

NULL AND false? Can only be false

left OR right: True if any one is true

NULL OR true? Must be true, no matter what value

NULL OR false? Could be true if NULL was true: = NULL

JOINS

```
SELECT [DISTINCT] target_list
FROM tableA, tableB
WHERE tableA.col = tableB.col AND ...
```

```
SELECT [DISTINCT] target_list
FROM tableA JOIN tableB
ON tableA.col = tableB.col
WHERE ...
```

(explicit) JOINS

```
SELECT [DISTINCT] target_list
FROM table_name
  [INNER {LEFT | RIGHT | FULL} {OUTER}] JOIN table_name
  ON qualification_list
WHERE ...
```

INNER is default

Difference is how to deal with NULL values

PostgreSQL documentation:
<http://www.postgresql.org/docs/9.4/static/tutorial-join.html>

Inner/Natural Join

```
SELECT s.sid, s.name, r.bid
FROM Sailors s, Reserves r
WHERE s.sid = r.sid
```

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

```
SELECT s.sid, s.name, r.bid
FROM Sailors s NATURAL JOIN Reserves r
```

All
Equivalent!

Natural Join means equi-join for each pair of
attrs with same name

Sailor names and their reserved boat ids

```

SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid

```

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

sid	bid	day
1	102	9/12
2	102	9/13

sid	name	bid
1	Eugene	102
2	Luis	102

Sailor names and their reserved boat ids

```

SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid

```

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

sid	bid	day
1	102	9/12
2	102	9/13

sid	name	bid
1	Eugene	102
2	Luis	102

Prefer INNER JOIN over NATURAL JOIN. Why?

Sailor names and their reserved boat ids

```

SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid

```

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

sid	bid	day
1	102	9/12
2	102	9/13

sid	name	bid
1	Eugene	102
2	Luis	102

Notice: No result for Ken!

Left Outer Join (or No Results for Ken)

Returns all matched rows *and all unmatched rows from table on left of join clause*
(at least one row for each row in left table)

```

SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid

```

All sailors & bid for boat in their reservations
Bid set to NULL if no reservation

Left Outer Join

```

SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid

```

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

sid	bid	day
1	102	9/12
2	102	9/13

sid	name	bid
1	Eugene	102
2	Luis	102
3	Ken	NULL

Can Left Outer Join be expressed with Cross-Product?

```

SELECT s.sid, s.name, r.bid
FROM Sailors s CROSS JOIN Reserves r
ON s.sid = r.sid

```

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

sid	bid	day
1	102	9/12
2	102	9/13

sid	name	bid
1	Eugene	102
2	Luis	102
3	Ken	102

sid	name	bid
1	Eugene	102
2	Luis	102
3	Ken	102

Right Outer Join

Same as LEFT OUTER JOIN, but guarantees result for rows in table on **right side of JOIN**

```
SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid
```

```
SELECT s.sid, s.name, r.bid
FROM Reserves r RIGHT OUTER JOIN Sailors s
ON s.sid = r.sid
```

RIGHT OUTER JOIN

```
SELECT s.sid, s.name, r.bid
FROM Sailors s RIGHT OUTER JOIN Reserves r
ON s.sid = r.sid
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13
4	109	9/20

Result

sid	name	bid
1	Eugene	102
2	Luis	102
NULL	NULL	109

Why is sid NULL?

FULL OUTER JOIN

Returns all matched or unmatched rows from both sides of JOIN

```
SELECT s.sid, s.name, r.bid
FROM Sailors s FULL OUTER JOIN Reserves r
ON s.sid = r.sid
```

FULL OUTER JOIN

```
SELECT s.sid, s.name, r.bid
FROM Sailors s Full OUTER JOIN Reserves r
ON s.sid = r.sid
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13
4	109	9/20

Result

sid	name	bid
1	Eugene	102
2	Luis	102
3	Ken	NULL
Left		
Right	NULL	109

JOIN Advice

Prefer "FROM TableA, TableB WHERE ..."
Except when you need OUTER JOIN (rare)

Integrity Constraints

Conditions that every legal instance must satisfy
Inserts/Deletes/Updates that violate ICs rejected
Helps ensure app semantics or prevent inconsistencies

We've discussed

domain/type constraints, primary/foreign key
general constraints ←

Beyond Keys: Table Constraints

Additional checks to ensure all data in table is valid

```
CREATE TABLE Sailors(
  sid int,
  ...
  PRIMARY KEY (sid),
  CHECK (rating >= 1 AND rating <= 10)
```

Nested subqueries
Named constraints

```
CREATE TABLE Reserves(
  sid int,
  bid int,
  day date,
  PRIMARY KEY (bid, day),
  CONSTRAINT no_red_reservations
  CHECK ('red' NOT IN (SELECT B.color
    FROM Boats B
    WHERE B.bid = bid))
```

WHAT!

So many things we can't express or don't work!

Nested queries in CHECK constraints



Advanced Stuff

User defined functions

Triggers

WITH

Views

User Defined Functions (UDFs)

Custom functions that can be called in database

Many languages: SQL, python, C, perl, etc

```
CREATE FUNCTION function_name(p1 type, p2 type, ...)
RETURNS type
```

User Defined Functions (UDFs)

Custom functions that can be called in database

Many languages: SQL, python, C, perl, etc

```
CREATE FUNCTION function_name(p1 type, p2 type, ...)
RETURNS type
AS $$
```

```
-- Logic
```

```
$$ LANGUAGE language_name;
```

User Defined Functions (UDFs)

Custom functions that can be called in database

Many languages: SQL, python, C, perl, etc

```
CREATE FUNCTION function_name(p1 type, p2 type, ...)
RETURNS type
AS $$
```

```
-- Logic
```

```
$$ LANGUAGE language_name;
```

A simple UDF (lang = SQL)

```
CREATE FUNCTION mult1(v int) RETURNS int
AS $$
SELECT v * 100;
$$ LANGUAGE SQL;
```

← Last statement is returned

```
CREATE FUNCTION function_name(p1 type, p2 type, ...)
RETURNS type
AS $$

-- Logic

$$ LANGUAGE language_name;
```

A simple UDF (lang = SQL)

```
CREATE FUNCTION mult1(v int) RETURNS int
AS $$
SELECT v * 100;
$$ LANGUAGE SQL;
```

```
SELECT mult1(S.age)
FROM sailors AS S
```

Sailors				Result
sid	name	rating	age	int4
1	Eugene	7	22	2200
2	Luis	2	39	3900
3	Ken	8	27	2700

<http://www.postgresql.org/docs/9.1/static/func-sql.html>

A simple UDF (lang = SQL)

```
CREATE FUNCTION mult1(int) RETURNS int
AS $$
SELECT $1 * 100;
$$ LANGUAGE SQL;
```

```
SELECT mult1(S.age)
FROM sailors AS S
```

Sailors				Result
sid	name	rating	age	int4
1	Eugene	7	22	2200
2	Luis	2	39	3900
3	Ken	8	27	2700

<http://www.postgresql.org/docs/9.1/static/func-sql.html>

Process a Record (lang = SQL)

```
CREATE FUNCTION mult2(x sailors) RETURNS int
AS $$
SELECT (x.sid + x.age) / x.rating;
$$ LANGUAGE SQL;
```

```
SELECT mult2(S.*)
FROM sailors AS S
```

Sailors				Result
sid	name	rating	age	int4
1	Eugene	7	22	3.285
2	Luis	2	39	20.5
3	Ken	8	27	3.75

<http://www.postgresql.org/docs/9.1/static/func-sql.html>

Process a Record (lang = SQL)

```
CREATE FUNCTION mult2(sailors) RETURNS int
AS $$
SELECT ($1.sid + $1.age) / $1.rating;
$$ LANGUAGE SQL;
```

```
SELECT mult2(S.*)
FROM sailors AS S
```

Sailors				Result
sid	name	rating	age	int4
1	Eugene	7	22	3.285
2	Luis	2	39	20.5
3	Ken	8	27	3.75

<http://www.postgresql.org/docs/9.1/static/func-sql.html>

Procedural Language/SQL (lang = plpgsql)

```
CREATE FUNCTION proc(v int) RETURNS int
AS $$
DECLARE
-- define variables
BEGIN
-- PL/SQL code
END;
$$ LANGUAGE plpgsql;
```

Boilerplate

<http://www.postgresql.org/docs/9.1/static/plpgsql.html>

Procedural Language/SQL (lang = plsql)

```
CREATE FUNCTION proc(v int) RETURNS int
AS $$
DECLARE
    -- define variables.  VAR TYPE [= value]
    qty int = 10;
BEGIN
    qty = qty * v;
    INSERT INTO blah VALUES(qty);
    RETURN qty + 2;
END;
$$ LANGUAGE plpgsql;
```

<http://www.postgresql.org/docs/9.4/static/plpgsql.html>

Procedural Code (lang = plpython2u)

```
CREATE FUNCTION proc(v int) RETURNS int
AS $$
import random
return random.randint(0, 100) * v
$$ LANGUAGE plpython2u;
```

Very powerful – can do anything so must be careful

run in a python interpreter with no security protection

plpy module provides database access

```
plpy.execute("select 1")
```

<http://www.postgresql.org/docs/9.4/static/plpython.html>

Procedural Code (lang = plpython2u)

```
CREATE FUNCTION proc(word text) RETURNS text
AS $$
import requests
resp = requests.get('http://google.com/search?q=%s' % v)
return resp.content.decode('unicode-escape')
$$ LANGUAGE plpython2u;
```

Very powerful – can do anything so must be careful

run in a python interpreter with no security protection

plpy module provides database access

```
plpy.execute("select 1")
```

<http://www.postgresql.org/docs/9.4/static/plpython.html>