

## Administrivia


Proj I Part 3 Due 11:59PM EST

DB Connection issues

Virtual env + pip install vs sudo pip install

## L20 Indexing Continued

### Alternative: Use a directory



Directory entries track #free bytes on data pages

Directory is collection of pages

## Indexes

“If I had eight hours to chop down a tree,  
I'd spend six sharpening my ax.”

*Abraham Lincoln*

## Indexes

Heap files can get data by sequential scan

Queries use *qualifications* (predicates)

find students in “CS”  
find students from CA

Indexes

file structures for value-based queries  
B+-tree index (~1970s)  
Hash index

Overview! Details in 4112

## Indexes

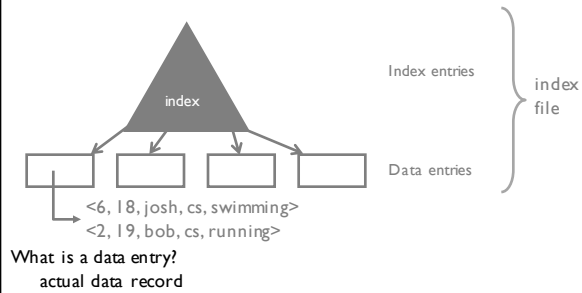
Defined wrt a *search key*  
different than a *candidate key*!

Faster access for WHERE clauses w/ search key

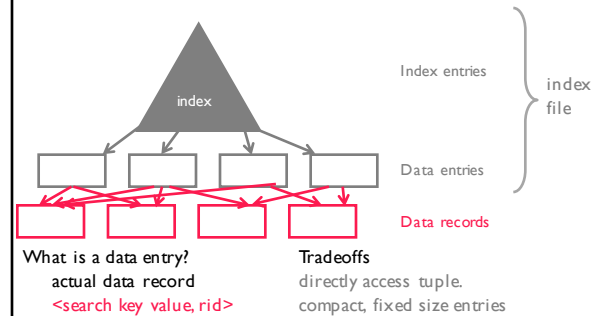
```
CREATE INDEX idx1 ON users USING btree (sid)
CREATE INDEX idx2 ON users USING hash (sid)
CREATE INDEX idx3 ON users USING btree (age,name)
```

You will play around with indexes in HW4

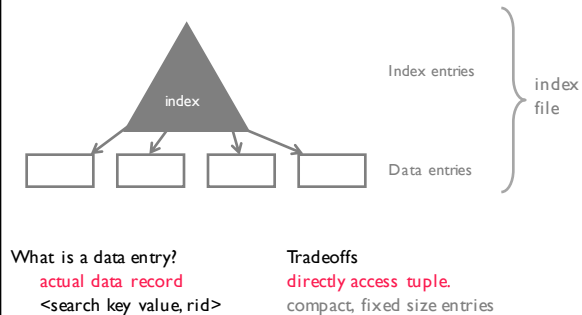
### High level (Primary) index structure



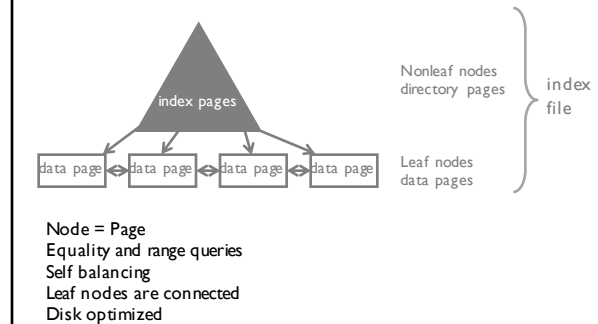
### High level (Secondary) index structure




### High level index structure



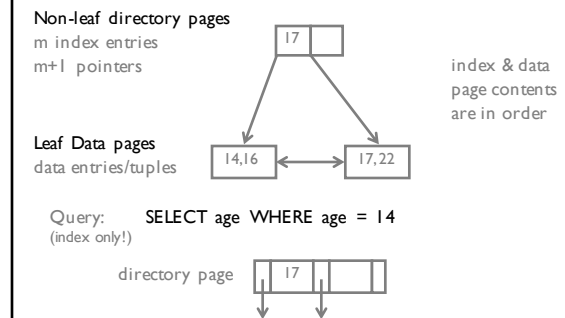
### B+ Tree Index




### B+ Tree on (age)



### Index Only Queries: B+ Tree on (age)




### B+ Tree on (age)



Query: `SELECT * WHERE age = 50`




### B+ Tree on (age)



### B+ Tree on (age)




### B+ Tree on (age)



Query: `SELECT * WHERE age > 15`

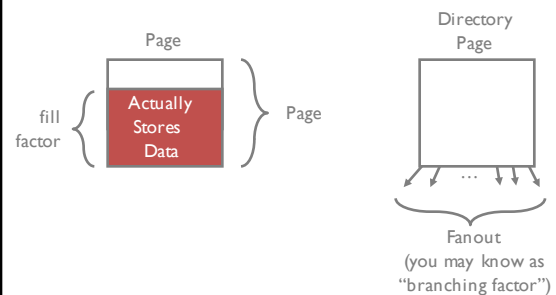
### B+ Tree on (age, name)



How do the following queries use the index on (age, name)?

`SELECT age WHERE age = 14`  
`SELECT * WHERE age < 18 AND name < 'monica'`  
`SELECT age WHERE name = 'bobby'`

### Terminology



## Some numbers (8kb pages)

How many levels?

fill-factor: ~66%

~300 entries per directory page

height 2:  $300^3 \sim 27$  Million entries

height 3:  $300^4 \sim 8.1$  Billion entries

Top levels often in memory

height 2 only 300 pages ~2.4MB

height 3 only 90k pages ~750MB

Cool B+Tree viz: <https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>

## Hash Index on age

Hash function

$$h(v) = v \% 3$$

Hash buckets  
containing  
data pages

0, 6

4, 13

5, 14

## INSERT Hash Index on age

Search key

1

Hash function

$$h(v) = v \% 3$$

Hash buckets  
containing  
data pages

0, 6

4, 13

5, 14

1

## INSERT Hash Index on age

Search key

11

Hash function

$$h(v) = v \% 3$$

Hash buckets  
containing  
data pages

0, 6

4, 13

5, 14

1

## INSERT Hash Index on age

Search key

11

Hash function

$$h(v) = v \% 3$$

Hash buckets  
containing  
data pages

0, 6

4, 13

5, 14

1

11

## SEARCH Hash Index on age

Search key

13

Hash function

$$h(v) = v \% 3$$

Hash buckets  
containing  
data pages

0, 6

4, 13

5, 14

1

11

Good for equality selections

Index = data pages + overflow data pages

Hash function  $h(v)$  takes as input the search key

## Costs

Three file types

Heap, B+ Tree, Hash

Operations we care about

Scan all data `SELECT * FROM R`

Equality `SELECT * FROM R WHERE x = I`

Range `SELECT * FROM R WHERE 10 < x and x < 50`

Insert record

Delete record

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything				
Equality				
Range				
Insert				
Delete				

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD			
Equality	0.5BD (avg)			
Range	BD			
Insert	2D			
Delete	Search + D			

Heap File

equality on a key. How many results?

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD		
Equality	0.5BD	$D(\log_2 B)$		
Range	BD	$D(\log_2 B + M)$		
Insert	2D	Search + BD		
Delete	Search + D	Search + BD		

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.25BD	
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} 1.25B + 1)$	
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} 1.25B + M)$	
Insert	2D	Search + BD	$D(\log_{80} 1.25B + 2)$	
Delete	Search + D	Search + BD	$D(\log_{80} 1.25B + 2)$	

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

B+ Tree

100 entries/directory page

80% fill factor

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.25BD	1.25BD
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} 1.25B + 1)$	D
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} 1.25B + M)$	1.25BD
Insert	2D	Search + BD	$D(\log_{80} 1.25B + 2)$	2D
Delete	Search + D	Search + BD	$D(\log_{80} 1.25B + 2)$	2D

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

B+ Tree

100 entries/directory page

80% fill factor

Hash index

no overflow

80% fill factor

**B** # data pages

**D** time to read/write page

**M** # pages in range query

## How to pick?

Depends on your queries (workload)

Which relations?

Which attributes?

Which types of predicates (=, <, >)

Selectivity

Insert/delete/update queries? how many?

## Naïve Algorithm

```
get query workload
group queries by type
for each query type in order of importance
    calculate best cost using current indexes
    if new index IDX will further reduce cost
        create IDX
```

Why not create every index?

updates are slower: upkeep costs  
takes up space

## High level guidelines

Check the WHERE clauses

attributes in WHERE are search/index keys

equality predicate → hash index

range predicate → tree index

Multi-attribute search keys supported

order of attributes matters for range queries

may enable queries that don't look at data pages (*index-only*)

## Summary

Design depends on economics, access cost ratios

Disk still dominant wrt cost/capacity ratio

Many physical layouts for files

same APIs, difference performance

remember physical independence

Indexes

Structures to speed up read queries

Multiple indexes possible

Decision depends on workload

## Things to Know

- How a hard drive works and its major performance characteristics
- The storage hierarchy and rough performance differences between RAM, SSD, Hard drives
- What files, pages, and records are, and how they are different than the UNIX model
- Heap File data structure
- B+ tree and Hash indexes
- Performance characteristics of different file organizations

## L20

## Query Execution & Optimization

## Steps for a New Application

Requirements  
what are you going to build?  
Conceptual Database Design  
pen-and-pencil description  
Logical Design  
formal database schema  
Schema Refinement  
fix potential problems, normalization  
Physical Database Design  
optimize for speed/storage  
App/Security Design  
prevent security problems

Optimization

## Recall

Relational algebra

equivalence: multiple stmts for same query  
some statements (much) faster than others

Which is faster?

- $\sigma_{v=1}(R \times T)$
- $\sigma_{v=1}(R) \times T$

$|R| = |T| = 10$  pages. 100? 1M?

# unique values in  $R = 1$ . 100? 1M? ← selectivity!


## Overview of Query Optimization

SQL  $\rightarrow$  query plan  
How plans are executed  
Some implementations of operators  
Cost estimation of a plan  
Selectivity  
System R dynamic programming

All ideas from System R's "Selinger Optimizer" 1979

## iPhones as a database

"avg acceleration over the past hour"



## SQL $\rightarrow$ Query Plan

SELECT a FROM R	$\pi_a(R)$	$\pi_a$   R
SELECT a FROM R WHERE a > 10	$\pi_a(\sigma_{a>10}(R))$	$\pi_a$   $\sigma_{a>10}$   R
SELECT a FROM R JOIN S ON R.b = S.b	$\pi_a(\bowtie_b(R,S))$	$\pi_a$   $\bowtie_b$ / \ R S