


Selinger Optimizer

parens(N) = 1
Only left-deep plans
ensures pipelining

Dynamic Programming

Idea: If considering ((ABC)DE)
compute best (ABC), cache, and reuse
figure out best way to combine with (DE)

Dynamic Programming Algorithm
compute best join size 1, then size 2, ...
 $\sim O(N^2)$



Summary

Single operator optimizations

Access paths
Primary vs secondary index costs
Projection/distinct
Predicate/project push downs

2 operators aka Joins

Nested loops, index nested loops, sort merge

Full plan optimizations

Naïve vs Selinger join ordering

Selectivity estimation

Statistics and simple models

Summary

Query optimization is a deep, complex topic
Pipelined plan execution
Different types of joins
Cost estimation of single and multiple operators
Join ordering is hard!

You should understand

Estimate query cardinality, selectivity
Apply predicate push down
Given primary/secondary indexes and statistics,
pick best index for access method + est cost
pick best index for join + est cost
pick cheaper of two execution plans
Given 3 or 4-way join
estimate join selectivity
pick the best join order

Transactions, Concurrency, Recovery

Transfer \$1000 from Evan to Neha

Check if Evan has \$1000
Evan's Account -= \$1000
Neha's Account += \$1000

Transfer \$1000 from Evan to Neha

Check if Evan has \$1000

Evan's Account -= \$1000

~~Neha's Account += \$1000~~ **Program crash or
user presses cancel:
Money disappeared**

Transfer \$1000 from Evan to Neha

Check if Evan has \$1000

~~Evan's Account -= \$1000~~

~~Neha's Account += \$1000~~

**OOPS! Not
enough money**

Two transfers: Starting with \$1500

Check if Evan has \$1000

Check if Evan has \$1000

Evan's Account -= \$1000

Evan's Account -= \$1000 **Negative
balance!**

Neha's Account += \$1000

Eugene's Account += \$1000

Transactions

Sequence of actions treated as a single unit

Atomicity: Apply all changes or none
("atomic" because it is indivisible)
Solves the crash problem

Isolation: Illusion that each transaction
executes sequentially, without concurrency

Transaction Guarantees

Atomicity

"all or nothing": All changes applied, or none are
users never see in-between transaction state

Consistency

database always satisfies Integrity Constraints
Transactions move from valid database to valid database

Isolation:

from transaction's point of view, it's the only one running

Durability:

if transaction commits, its effects *must persist*

Transactions

Transaction: a sequence of actions

action = read object, write object, commit, abort
API between app semantics and DBMS's view

User's view

T1: begin A=A+100 B=B-100 END

T2: begin A=A-50 B=B+50 END

DBMS's logical view

T1: begin r(A) w(A) r(B) w(B) END

T2: begin r(A) w(A) r(B) w(B) END

Concepts

Concurrency Control

techniques to ensure **correct** results when running transactions concurrently

what does this mean?

Recovery

On crash or abort, how to get back to a consistent (**correct**) state?

The two are intertwined!

What is Correct?

Serializability

Regardless of the interleaving of operations, result same as a serial ordering

Schedule

One specific interleaving of the operations

Serial Schedules

Logical xacts

T1: r(A) w(A) r(B) w(B) (e.g. A=A+100; B=B-100)
T2: r(A) w(A) r(B) w(B) (e.g. A=A*1.5; B=B*1.5)

No concurrency (**serial 1**)

T1: r(A) w(A) r(B) w(B)
T2: r(A) w(A) r(B) w(B)

No concurrency (**serial 2**)

T1: r(A) w(A) r(B) w(B)
T2: r(A) w(A) r(B) w(B)

Are serial 1 and serial 2 equivalent?

More Example Schedules

Logical xacts

e.g. A=0 B=0

T1: r(A) w(A) **r(A)** w(B) e.g. A=A+1; B=A+1
T2: r(A) w(A) r(B) w(B) e.g. A=A+10; B=B+1

Concurrency (bad)

T1: r(A) w(A) r(A) w(B)
T2: r(A) w(A) r(B) w(B)

Concurrency (same as serial T1, T2!)

T1: r(A) w(A) r(A) w(B)
T2: r(A) w(A) w(A) r(B) w(B)

Concepts

Serial schedule

One transaction at a time. no concurrency.

Equivalent schedule

the database state is the same at end of both schedules

Serializable schedule (gold standard)

equivalent to a **serial** schedule

SQL → R/W Operations

```
UPDATE accounts
SET bal = bal + 1000
WHERE bal > 1M
```

Read all balances for every tuple

Update those with balances > 1M

Does the access method matter?

Why Serializable Schedule? Anomalies

Reading in-between (uncommitted) data

T1: R(A) W(A) R(B) W(B) abort
T2: R(A) W(A) commit
WR conflict or dirty reads

Reading same data gets different values

T1: R(A) R(A) W(A) commit
T2: R(A) W(A) commit
RW conflict or unrepeatable reads

Why Serializable Schedule? Anomalies

Stepping on someone else's writes

T1: W(A) W(B) commit
T2: W(A) W(B) commit
WW conflict or lost writes

Notice: all anomalies involve writing to data that is read/written to.

If we track our writes, maybe can prevent anomalies

Conflict Serializability

What is a conflict?

For 2 operations, if run in different order, get different results

Conflict?	R	W
R	NO	YES
W	YES	YES

Conflict Serializability

def: a schedule that is conflict equivalent to a serial schedule

Meaning: you can swap non-conflicting operations to derive a serial schedule.

∀ conflicting operations O1 of T1, O2 of T2

O1 always before O2 in the schedule or

O2 always before O1 in the schedule