

Statistics

Median: 93
Mean: 90
Stddev: 15.9

Everything else about Databases

One Size Fits All



Consistent tools to access data
Widely understood
Rich, sophisticated queries, tools and features
Reliable

One Size **Does Not** Fits All



One Size **Does Not** Fits All



DBMSes in the Wild

Classic Relational

\$\$: Oracle, IBM, Microsoft, Teradata, EMC, etc
Free: MySQL, PostgreSQL

New Relational

In-Memory, Column-store, Streaming

Non-traditional

Search (Google, Bing, Lucene), Scientific, Geographic

NoSQL

Big Data: Hadoop, Spark, ...
Key-value: Mongo, Cassandra, Memcache, Redis, ...

DBMS-as-a-Service

Microsoft Azure, Amazon Redshift/RDS, etc...

DBMSes in the Wild

Classic Relational

\$\$: Oracle, IBM, Microsoft, Teradata, EMC, etc

Free: MySQL, PostgreSQL

New Relational

In-Memory, Column-store, Streaming

Non-traditional

Search (Google, Bing, Lucene), Scientific, Geographic

NoSQL

Big Data: Hadoop, Spark, etc

Key-value: Mongo, Cassandra, Memcache, Redis, ...

DBMS-as-a-Service

Microsoft Azure, Amazon Redshift/RDS, etc...

Modern Database Systems

90s: The Internet:

Every application is 24x7x365

Some applications have huge numbers of users

Traditional solutions fall over

Slashdot effect, Twitter fail whale, etc etc

Solution? Sharding (partitioning)


Split one database into many

Don't access different "shards" at once

Ebay: Shard per auction

Gmail: Shard per email

Many successes: eBay, Facebook, YouTube, Salesforce



Sharding challenges

Limits supported operations

e.g. No joins/transactions between shards

Hot/large/imbalanced shards

e.g. Huge customers, popular pages

Manage many database instances

Transfers many challenges back to application

Google Bigtable; 2006

Sharding means you can't use relational model

Use simpler model: Key/value

Store unique keys, associated with values

key: "evan"

Value: "adjunct:w4ll:ej@evanjones.ca"

Google Bigtable; 2006

Simple key/value model can be easily scaled

Split tables into tablets

Distribute tablets to multiple servers

Split tablets that have grown too big

Created the “NoSQL” movement

Other NoSQL systems

MongoDB, Cassandra: Distributed systems

Memcached (2003): Simple key/value in memory

Redis (2009): Key/value with lots of features!

Google Spanner, 2012

Globally distributed transactions

SQL interface

... basically a global relational database

Distributed databases today

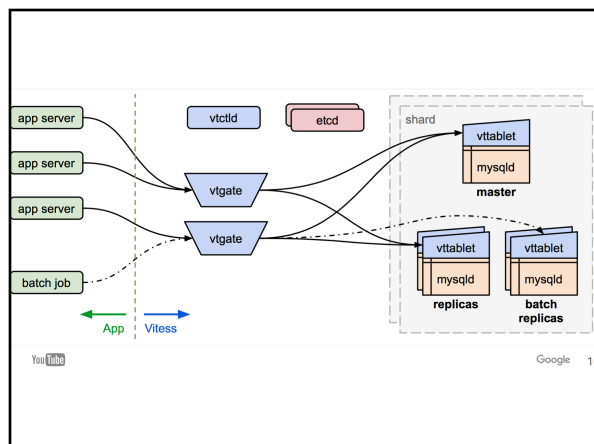
Surprisingly few that are widely used

Many analytic database product (next slides)

NoSQL: Cassandra, MongoDB, HBase

Many startups, no clear winners

Lots of sharded MySQL/Postgres/etc with custom tools: e.g. YouTube Vitess



Why no commercial products? Hypothesis:

A “commodity” server today is big:

32 CPUs, 208 GB RAM, \$1000/month cloud

Big enough for many apps

Extremely large apps: have own dev team

Sharding is painful, but does work

OLTP vs OLAP

OnLine Transaction Processing

Interactive queries, low latency
Small amount of data per transaction
Modifies data

OnLine Analytical Processing

Batch queries; "high" latency
Aggregates, summaries
Mostly read only


Data Warehouses

Store all historical data for future analysis

Sales by month over past 20 years
Clicks by youth in Texas
Cost by product component


Most companies have something that serves this purpose

Star Schema





Star Schema


Fact table is "fat"; Dimensions are denormalized
Queries access ~6 attrs



100 attributes




16x less data read. Unfair advantage.



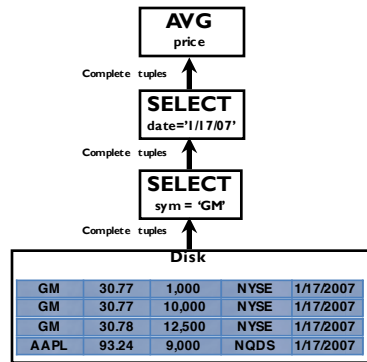
16x less data read. Unfair advantage.
Compression better on single column

Execute on compressed data



Traditional DBMS

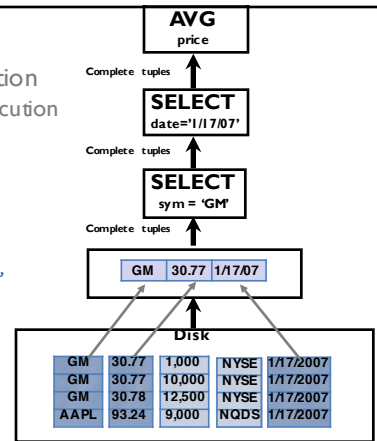
```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```



27

Naïve:
Early Materialization
Row oriented execution

```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```

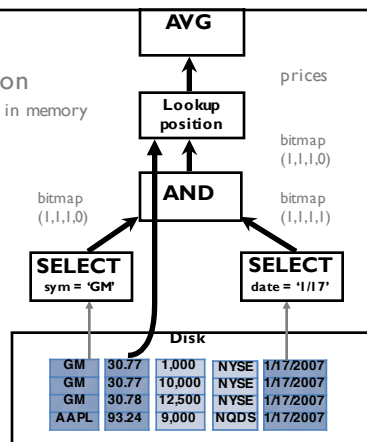


28

C-Store
Late Materialization

Much less data moving in memory

```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```

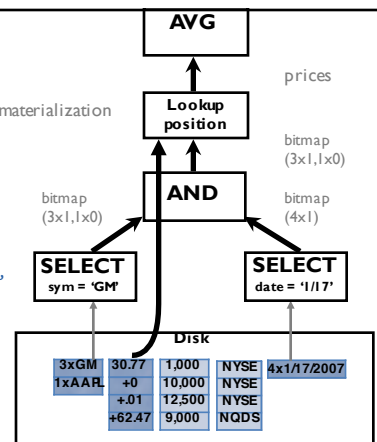


29

C-Store
Compression

Only possible w/ late materialization

```
SELECT avg(price)
FROM tickstore
WHERE symbol = 'GM'
AND date = '1/17/2007'
```



30

Column Stores

Optimized for data warehouses
 Store data by attribute/column rather than row
 Compression
 Compressed *query plan execution*

50-100x faster than row store

Column Stores

Optimized for data warehouses
 Store data by attribute/column rather than row
 Compression
 Compressed *query plan execution*

50-100x faster than row store
 (for OLAP queries)

Many successful products

HP Vertica
 Teradata
 Netezza

 ... many others

Meanwhile at the Internet companies

Record everything!
 On hundreds or thousands of computers!

 ... how do we do anything with it?

Google MapReduce 2004

map(records) \rightarrow (key,value)
 sort all keys
 reduce(key,[value]) \rightarrow (key2,value2)

 distributed to thousands of machines

Example: requests per day

map(request log record) \rightarrow (day, 1)
 sort all keys
 reduce(day,[1, 1,...]) \rightarrow (day, sum(counts))

Apache Hadoop, January 2006

Open source clone started by Doug Cutting
Cutting was at Yahoo!, working on search

Gained significant adoption within ~2 years
Cloudera started to commercialize
Hortonworks spun out of Yahoo much later

Hadoop versus Databases “debate”

Database industry: MapReduce is a bad implementation of distributed databases

MapReduce crowd: databases can't scale

More rational perspective

MapReduce:
Relatively easy to scale
Amazing for unstructured data
Manual work for every “query”

Databases:
Setup and “loading” takes work/time
Rich queries without much effort

Today:

Google Powerdrill/BigQuery: SQL interface
Cloudera Impala: SQL interface
Facebook Presto: SQL interface

Reason? Not perfect but well understood
Queries better than implementing joins by hand!
Optimizers frequently faster than humans

In-Memory DBMSes

Transaction-oriented apps
remove 1 unit from product
move 5 units from org 1 to org 2
(shopping carts, inventory)

Data stored in memory
Disk only used for recovery
Active-active replication for fault-tolerance



Traditional Database

Indexes	queries go faster
Concurrency	queries go faster
Locking	serializability (go slower)
Logging	recovery (go slower)
Buffer Manager	manage pages in memory (go slower)

make up your mind!

Results after removing the components (in # instruction)

Instruction of useful work is only <2% of a memory resident DB



In memory DB

Barebones
Executor

What about
Parsing
Concurrency?
Recovery?

In memory DB

Barebones
Executor

Stored Procedures

What about
Parsing
Concurrency?
Recovery?

Procedure:

```
p1 = SELECT cost
      FROM fact_table
      WHERE id = ?
```

Query:


```
p1(10)
```

In memory DB

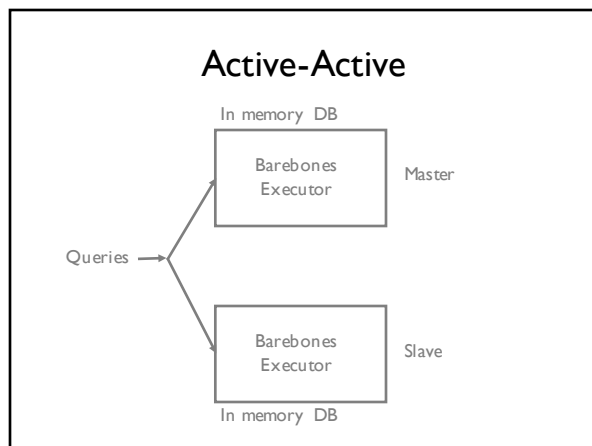
Barebones
Executor

What about
Parsing
~~Concurrency?~~
no buffer manager, no concurrency, no locks
Recovery?

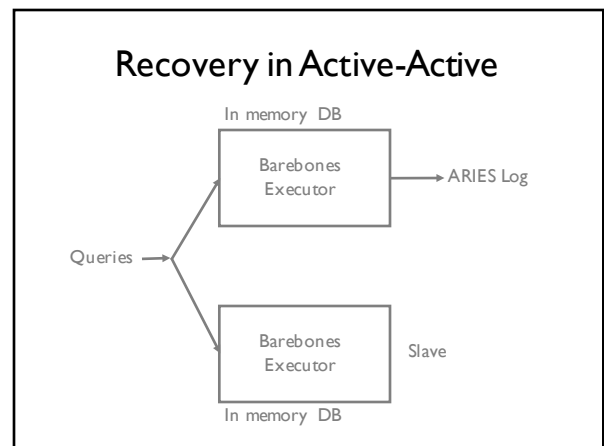
Log Shipping



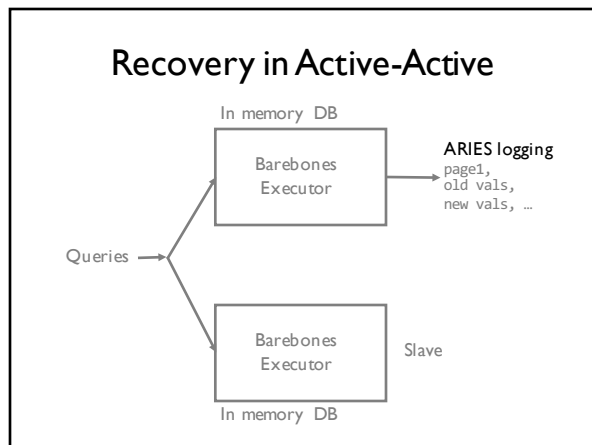
Active-Active



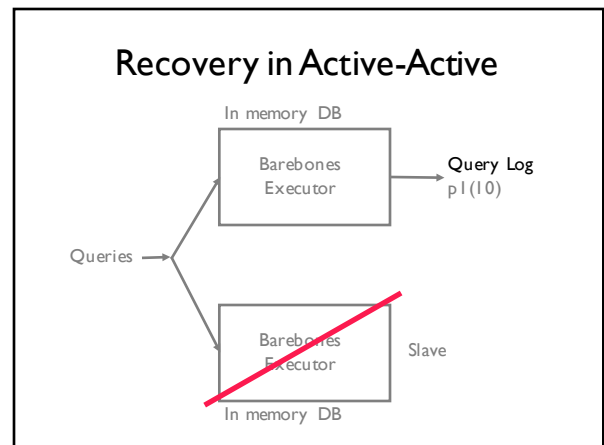
Recovery in Active-Active



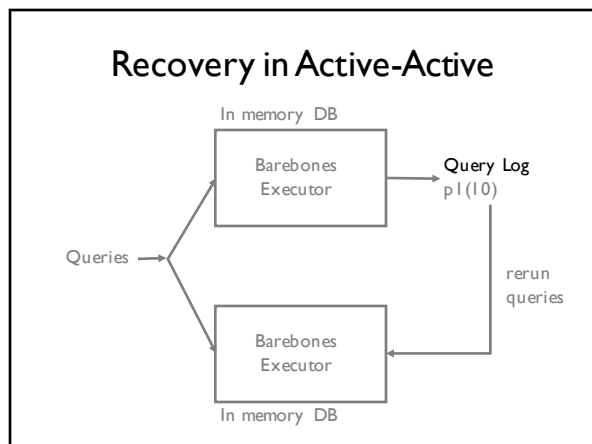
Recovery in Active-Active



Recovery in Active-Active



Recovery in Active-Active



Databases as a Service

Amazon:

Redshift (columnar)
RDS (traditional)
DynamoDB (key/value)

Google:

BigTable (key/value)
Cloud SQL (traditional)
BigQuery (columnar)

Microsoft:

SQL Azure
DocumentDB (key/val)

Database Service Challenges

Legal rules:

- Data residency (EU, others)
- Privacy/Compliance (e.g. HIPAA for health)
- Security/Compliance (e.g. PCI for credit cards)

Data gravity:

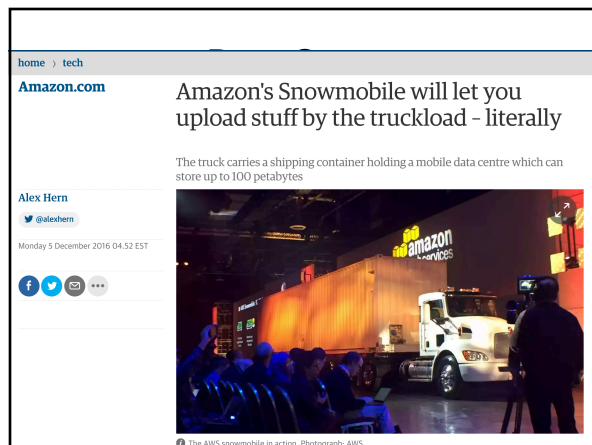
- Large data is hard to move
- Latency across the Internet can be bad

Dave Patterson interviews Jim Gray

DP "Sneaker net" was when you used your sneakers to transport data?

JG In the old days, sneaker net was the notion that you would pull out floppy disks, run across the room in your sneakers, and plug the floppy into another machine. This is just TeraScale SneakerNet. You write your terabytes onto this thing and ship it out to your pals. Some of our pals are extremely well connected—they are part of Internet 2, Virtual Business Networks (VBNs), and the Next Generation Internet (NGI). Even so, it takes them a long time to copy a gigabyte. Copy a terabyte? It takes them a very, very long time across the networks they have.


<http://queue.acm.org/detail.cfm?id=864078>



One Size Does Not Fits All



"Modern" Systems: Connected!



Other Aspects of Database Research

- Domain specific data management
- Data quality/cleaning
- Database usability
- Crowdsourced Databases
- Information extraction and mining
- Query compilation
- Data visualization and exploration

Societal Data

Health	Fake data
Investigative Journalism	Biased data
Recommendations	Incorrect data
Politics	Mixed data
Surveillance	
Identity	

Data will be crucial to
how we live
as individuals and as a society

Go forth and make a
(positive) difference!